CEREBRO Ground Truth Benchmark (CEREBRO-GT v0.1)

Benchmark · CEREBRO · Joselito Dizon · 2026-09-05

A frozen, annotated benchmark for measuring the knowledge pipeline: expected fragments, concepts and provenance for each source.

Why a benchmark

The Practice Module Brief rewards evaluation against ground truth, so the pipeline needs a fixed reference before it is tuned.

This benchmark evaluates the pipeline described in the CEREBRO Solution Architecture.

Annotation schema

Each benchmark item records the source artifact, its checksum, the expected fragments with character ranges, and the expected concepts.

Annotations are frozen as version 0.1 and persisted with their own checksum, so later changes are visible.

Acceptance tests

A pipeline run passes when every expected fragment is produced with the right source location and no fragment lacks provenance.

Concept extraction is scored separately, because wording can differ while meaning is preserved.