Local LLM Evaluation Plan

Plan · Infrastructure · Joselito Dizon · 2026-09-28

How small local models will be compared for grounded answering in both projects; measurements are still to be run on Buddy.

Goal

Choose a local model for answer phrasing in the MealSenseAI Meal Advisor and for the reasoning stage in the CEREBRO Solution Architecture.

No external LLM API is used; every candidate runs on Buddy through Ollama.

Method

Each candidate phrases answers over a labelled set; a validator counts unsupported numbers and entities, following the Knowledge Graphs and RAG lecture notes.

Latency, tokens per second and GPU memory will be measured on the server, not estimated.

Status

The labelled sets exist; the measurements on Buddy have not been run yet, so no results are reported here.