CEREBRO Ground Truth Benchmark (CEREBRO-GT v0.1)

Prepared by Joselito Dizon

A frozen, annotated benchmark for measuring the knowledge pipeline before it is tuned.

Why a benchmark

The benchmark was frozen on 5 September 2026, before any tuning. Model Evaluation without a fixed reference only measures how the pipeline changed, not whether it is right.

Annotation schema

Each item records the source artefact and its checksum, the expected fragments with their character ranges, and the expected concepts. Annotation itself is Human-in-the-Loop Review: a person approves every label.

Acceptance tests

A pipeline run passes when every expected fragment is produced with the right Provenance and no fragment lacks a source location. Concept extraction is scored separately, because wording can differ while meaning holds.

Illustrative sample document for the CEREBRO demo library.