Local LLM Evaluation Plan

Prepared by Joselito Dizon

How small local models will be compared for grounded answering in both projects.

Goal

Planned on 28 September 2026. Choose Local Language Models for answer phrasing in the MealSenseAI Meal Advisor and for Retrieval-Augmented Generation in CEREBRO Recall. Every candidate runs through Ollama; no cloud API is used.

Candidates

qwen3 for answering, because it follows citation instructions well once its deliberation is stripped; mxbai-embed-large for embeddings, because its vectors separate topics cleanly in a small library.

Method

Model Evaluation uses a labelled question set. A validator counts unsupported numbers and named things in each answer; latency, tokens per second and GPU memory are measured on the home server, not estimated.

Illustrative sample document for the CEREBRO demo library.